

m91br7gdw

February 17, 2026

Task 02: Credit Risk Prediction.

Objective: Predict whether a loan applicant is likely to default on a loan.

Dataset: Loan Prediction Dataset(available on Kaggle). About Dataset: 1. Borrower's Demographics: age (int64) – Borrower's age (in years). gender (category) – Borrower's gender (Male/Female). marital_status (category) – Marital status (Single, Married, Divorced). education_level (category) – Education level (High School, Bachelor, Master, PhD). 2. Financial Information: annual_income (float64) – Borrower's yearly income. monthly_income (float64) – Borrower's monthly income. employment_status (category) – Current employment type (Employed, Self-Employed, Unemployed). debt_to_income_ratio (float64) – Ratio of borrower's debt to their income. Lower = better. credit_score (int64) – Credit bureau score (e.g., FICO). Higher = less risky. 3. Loan Information: loan_amount (float64) – Amount of loan taken. loan_purpose (category) – Loan purpose (Car, Education, Home, Medical, etc.). interest_rate (float64) – Loan par annual interest rate (%). loan_term (int64) – Loan repayment duration (months, e.g., 36 or 60). installment (float64) – Monthly installment . grade_subgrade (category) – Risk category assigned to loan (A1, B2, etc.). 4. Borrower's Credit History: num_of_open_accounts (int64) – Total active credit accounts. total_credit_limit (float64) – Borrower's total available credit limit. current_balance (float64) – Borrower's outstanding balance (loan + credit card). delinquency_history (int64) – Count of late payments in borrower's history. public_records (int64) – Negative public records (e.g., bankruptcies, legal actions). num_of_delinquencies (int64) – Total delinquencies (missed payments). 5. Target Variable: loan_paid_back (int64) – Target variable: 1 → Borrower paid loan in full. 0 → Borrower defaulted (did not repay fully).

Loading the Dataset.

```
[1]: import pandas as pd
```

```
[2]: Dataset = pd.read_csv("loan_dataset_20000.csv")
```

Data Exploration.

```
[3]: Dataset.head()
```

```
[3]:   age  gender  marital_status  education_level  annual_income  monthly_income \
0    59     Male        Married      Master's       24240.19      2020.02
1    72   Female        Married    Bachelor's       20172.98      1681.08
2    49   Female        Single    High School       26181.80      2181.82
3    35   Female        Single    High School       11873.84      989.49
4    63    Other        Single         Other       25326.44      2110.54
```

```

employment_status debt_to_income_ratio credit_score loan_amount ... \
0      Employed          0.074        743    17173.72 ...
1      Employed          0.219        531    22663.89 ...
2      Employed          0.234        779    3631.36 ...
3      Employed          0.264        809   14939.23 ...
4      Employed          0.260        663   16551.71 ...

loan_term installment grade_subgrade num_of_open_accounts \
0       36      581.88           B5             7
1       60      573.17           F1             5
2       60      76.32            B4             2
3       36      468.07           A5             7
4       60      395.50           D5             1

total_credit_limit current_balance delinquency_history public_records \
0      40833.47      24302.07             1             0
1      27968.01      10803.01             1             0
2      15502.25      4505.44              0             0
3      18157.79      5525.63              4             0
4      17467.56      3593.91              2             0

num_of_delinquencies loan_paid_back
0                  1                 1
1                  3                 1
2                  0                 1
3                  5                 1
4                  2                 1

```

[5 rows x 22 columns]

[4]: `Dataset.tail()`

```

[4]:      age gender marital_status education_level annual_income \
19995    39 Female     Married Bachelor's    39640.08
19996    66 Female     Married Bachelor's    32062.90
19997    65 Female     Single  Master's     18642.02
19998    35 Male       Married Master's     22181.39
19999    36 Female     Married Other        23737.70

monthly_income employment_status debt_to_income_ratio credit_score \
19995      3303.34           Employed          0.275        691
19996      2671.91           Employed          0.367        758
19997      1553.50           Student          0.106        751
19998      1848.45           Retired          0.275        646
19999      1978.14           Employed          0.228        630

```

```

    loan_amount ... loan_term installment grade_subgrade \
19995      16322.23 ...       36      566.22           C5
19996      16697.34 ...       36      553.71           B5
19997      23924.78 ...       36      772.66           B4
19998      16920.13 ...       36      595.36           D2
19999      15769.75 ...       36      531.88           D2

    num_of_open_accounts total_credit_limit current_balance \
19995                  2            23748.10      5801.45
19996                  8            49929.65     40901.31
19997                  3            13137.57      5075.67
19998                  5            19580.82      3876.16
19999                  8            43013.59     12753.03

    delinquency_history public_records num_of_delinquencies \
19995                 1             0               4
19996                 3             0               3
19997                 1             0               2
19998                 4             0               5
19999                 2             0               2

    loan_paid_back
19995          0
19996          1
19997          1
19998          1
19999          0

```

[5 rows x 22 columns]

[5]: Dataset.shape

[5]: (20000, 22)

[6]: Dataset.info()

```

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 20000 entries, 0 to 19999
Data columns (total 22 columns):
 #   Column           Non-Null Count  Dtype  
--- 
 0   age              20000 non-null   int64  
 1   gender            20000 non-null   object  
 2   marital_status   20000 non-null   object  
 3   education_level  20000 non-null   object  
 4   annual_income    20000 non-null   float64 
 5   monthly_income   20000 non-null   float64 

```

```

6   employment_status      20000 non-null  object
7   debt_to_income_ratio  20000 non-null  float64
8   credit_score           20000 non-null  int64
9   loan_amount             20000 non-null  float64
10  loan_purpose            20000 non-null  object
11  interest_rate           20000 non-null  float64
12  loan_term                20000 non-null  int64
13  installment               20000 non-null  float64
14  grade_subgrade            20000 non-null  object
15  num_of_open_accounts     20000 non-null  int64
16  total_credit_limit       20000 non-null  float64
17  current_balance           20000 non-null  float64
18  delinquency_history      20000 non-null  int64
19  public_records              20000 non-null  int64
20  num_of_delinquencies      20000 non-null  int64
21  loan_paid_back              20000 non-null  int64
dtypes: float64(8), int64(8), object(6)
memory usage: 3.4+ MB

```

[7]: `Dataset.describe()`

	age	annual_income	monthly_income	debt_to_income_ratio	\
count	20000.000000	20000.000000	20000.000000	20000.000000	
mean	48.027000	43549.637765	3629.136466	0.177019	
std	15.829352	28668.579671	2389.048326	0.105059	
min	21.000000	6000.000000	500.000000	0.010000	
25%	35.000000	24260.752500	2021.730000	0.096000	
50%	48.000000	36585.260000	3048.770000	0.160000	
75%	62.000000	54677.917500	4556.495000	0.241000	
max	75.000000	400000.000000	33333.330000	0.667000	
	credit_score	loan_amount	interest_rate	loan_term	installment \
count	20000.000000	20000.000000	20000.000000	20000.000000	20000.000000
mean	679.25695	15129.300909	12.400627	43.22280	455.625794
std	69.63858	8605.405513	2.442729	11.00838	274.622125
min	373.00000	500.000000	3.140000	36.00000	9.430000
25%	632.00000	8852.695000	10.740000	36.00000	253.910000
50%	680.00000	14946.170000	12.400000	36.00000	435.595000
75%	727.00000	20998.867500	14.002500	60.00000	633.595000
max	850.00000	49039.690000	22.510000	60.00000	1685.400000
	num_of_open_accounts	total_credit_limit	current_balance	\	
count	20000.000000	20000.000000	20000.000000	20000.000000	
mean	5.011800	48649.824769	24333.394631		
std	2.244529	32423.378128	22313.845395		
min	0.000000	6157.800000	496.350000		
25%	3.000000	27180.492500	9592.572500		

```

50%           5.000000    40241.615000   18334.555000
75%           6.000000    60361.257500   31743.327500
max          15.000000   454394.190000  352177.900000

      delinquency_history  public_records  num_of_delinquencies \
count      20000.000000    20000.000000     20000.000000
mean       1.990150        0.061800      2.489150
std        1.474945        0.285105      1.631384
min        0.000000        0.000000      0.000000
25%        1.000000        0.000000      1.000000
50%        2.000000        0.000000      2.000000
75%        3.000000        0.000000      3.000000
max       11.000000        2.000000     11.000000

      loan_paid_back
count      20000.000000
mean       0.799900
std        0.400085
min        0.000000
25%        1.000000
50%        1.000000
75%        1.000000
max       1.000000

```

[8]: `Dataset.columns`

```

[8]: Index(['age', 'gender', 'marital_status', 'education_level', 'annual_income',
       'monthly_income', 'employment_status', 'debt_to_income_ratio',
       'credit_score', 'loan_amount', 'loan_purpose', 'interest_rate',
       'loan_term', 'installment', 'grade_subgrade', 'num_of_open_accounts',
       'total_credit_limit', 'current_balance', 'delinquency_history',
       'public_records', 'num_of_delinquencies', 'loan_paid_back'],
       dtype='object')

```

[9]: `Dataset.dtypes`

```

[9]: age            int64
      gender         object
      marital_status  object
      education_level object
      annual_income   float64
      monthly_income  float64
      employment_status object
      debt_to_income_ratio float64
      credit_score    int64
      loan_amount     float64
      loan_purpose    object

```

```
interest_rate          float64
loan_term              int64
installment            float64
grade_subgrade         object
num_of_open_accounts   int64
total_credit_limit    float64
current_balance        float64
delinquency_history   int64
public_records         int64
num_of_delinquencies  int64
loan_paid_back         int64
dtype: object
```

Visualize key features such as loan amount, education and income.

```
[10]: import matplotlib as pyplot
from matplotlib import pyplot as plt


import seaborn as sns
```

```
[11]: plt.figure(figsize=(6,6))

plt.hist(Dataset["loan_amount"])

plt.title("Distribution of Loan Amount")
plt.xlabel("Loan Amount")
plt.ylabel("Frequency of Applicants")

plt.show()
```




```
[12]: plt.figure(figsize=(6,6))

plt.hist(Dataset["education_level"])

plt.title("Distribution of Applicants by Education Level")
plt.xlabel("Education Level")
plt.ylabel("Frequency of Applicants")

plt.show()
```




```
[13]: plt.figure(figsize=(6,6))

plt.scatter(Dataset["education_level"], Dataset["loan_amount"])

plt.title("Education Level & Loan Amount")
plt.xlabel("Education Level")
plt.ylabel("Loan Amount")

plt.show()
```




```
[14]: plt.figure(figsize=(6,6))

plt.boxplot(Dataset["monthly_income"])

plt.title("Distribution of Monthly Income")
plt.xlabel("Monthly Income")
plt.ylabel("Frequency of Applicants")

plt.show()
```




```
[15]: plt.figure(figsize=(6,6))

plt.boxplot(Dataset["annual_income"])

plt.title("Distribution of Annual Income")
plt.xlabel("Annual Income")
plt.ylabel("Frequency of Applicants")

plt.show()
```




```
[16]: plt.figure(figsize=(6,6))

plt.scatter(Dataset["monthly_income"], Dataset["annual_income"])

plt.title("Monthly Income & Annual Income")
plt.xlabel("Monthly Income")
plt.ylabel("Annual Income")

plt.show()
```




```
[17]: plt.figure(figsize=(6,6))

sns.violinplot(x="loan_paid_back", y="monthly_income", data=Dataset)

plt.title("Monthly Income & Loan Paid Back")
plt.xlabel("Loan Paid Back")
plt.ylabel("Monthly Income")

plt.show()
```




```
[18]: plt.figure(figsize=(5,30))

sns.pairplot(Dataset)

plt.suptitle("Showing the relationships between all the variables.")

plt.show()
```

<Figure size 500x3000 with 0 Axes>



Instructions: Handle missing values appropriately.

[19]: # Import Libraries.

```

import pandas as pd
import numpy as np

import matplotlib.pyplot
from matplotlib import pyplot as plt
%matplotlib inline

import seaborn as sns

```

```
from sklearn.model_selection import train_test_split
from sklearn.impute import SimpleImputer
from sklearn.linear_model import LogisticRegression
from sklearn.tree import DecisionTreeClassifier

from sklearn.metrics import accuracy_score, confusion_matrix
```

[20]: # Check missing values per column.

```
print(Dataset.isnull().sum())

# This output means:
# No NaN values.
# No empty cells.
# No imputation required.
# Dataset is already clean for missing data.
```

```
age                      0
gender                   0
marital_status           0
education_level          0
annual_income             0
monthly_income            0
employment_status         0
debt_to_income_ratio     0
credit_score              0
loan_amount               0
loan_purpose              0
interest_rate              0
loan_term                  0
installment                 0
grade_subgrade             0
num_of_open_accounts       0
total_credit_limit         0
current_balance             0
delinquency_history        0
public_records              0
num_of_delinquencies        0
loan_paid_back              0
dtype: int64
```

[21]: # Dataset have no missing values. If Dataset have missing values then we use use ↵ the following method.

```
# Decide strategy:
# Numerical: Fill with median.
# Categorical: Fill with most frequent.
```

```

numerical_columns = Dataset.select_dtypes(include=["int64","float64"]).columns
categorical_columns = Dataset.select_dtypes(include=["object"]).columns

numerical_imputer = SimpleImputer(strategy="median")
categorical_imputer = SimpleImputer(strategy="most_frequent")

Dataset[numerical_columns] = numerical_imputer.
    ↪fit_transform(Dataset[numerical_columns])
Dataset[categorical_columns] = categorical_imputer.
    ↪fit_transform(Dataset[categorical_columns])

```

[22]: `print(Dataset.isnull().sum().sum())`

0

Train a Classification Model like Logistic Regression and Decision Tree.

Logistic Model 01.

[23]: `x = Dataset.drop("loan_paid_back", axis=1)`
`y = Dataset["loan_paid_back"]`

[24]: `x = pd.get_dummies(x, drop_first=True)`

[25]: `x_train, x_test, y_train, y_test = train_test_split(x, y, test_size=0.2, ↪random_state=42)`

[26]: `print("Training Data:", x_train.shape)`
`print("Testing Data:", x_test.shape)`

Training Data: (16000, 64)

Testing Data: (4000, 64)

[27]: `logistic_model = LogisticRegression(max_iter=100)`
`logistic_model.fit(x_train, y_train)`

`y_predict_logistic = logistic_model.predict(x_test)`

```
c:\Users\zaigh\AppData\Local\Programs\Python\Python313\Lib\site-
packages\sklearn\linear_model\_logistic.py:469: ConvergenceWarning: lbfgs failed
to converge (status=1):
STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.
```

Increase the number of iterations (max_iter) or scale the data as shown in:
<https://scikit-learn.org/stable/modules/preprocessing.html>

Please also refer to the documentation for alternative solver options:
https://scikit-learn.org/stable/modules/linear_model.html#logistic-

```
regression
n_iter_i = _check_optimize_result(
```

DecisionTree Model.

```
[28]: decisionTree_model = DecisionTreeClassifier(random_state=42)

decisionTree_model.fit(x_train, y_train)

y_predict_decisionTree = decisionTree_model.predict(x_test)
```

Evaluate the model using accuracy and a confusion matrix.

```
[29]: print("Logistic Regression Accuracy: \n",
          accuracy_score(y_test, y_predict_logistic))

print("Confusion Matrix: \n",
      confusion_matrix(y_test, y_predict_logistic))
```

Logistic Regression Accuracy:

0.7955

Confusion Matrix:

```
[[ 0  818]
 [ 0 3182]]
```

```
[30]: cm = confusion_matrix(y_test, y_predict_logistic)
TN, FP, FN, TP = cm.ravel()

# Create a Structured Table.
metrics_table = pd.DataFrame({
    'Metric': ['TP', 'TN', 'FP', 'FN', 'Accuracy', 'Precision', 'Recall'],
    'Value': [TP, TN, FP, FN,
              (TP+TN)/(TP+TN+FP+FN),
              TP/(TP+FP),
              TP/(TP+FN)]})

print(metrics_table)
```

	Metric	Value
0	TP	3182.0000
1	TN	0.0000
2	FP	818.0000
3	FN	0.0000
4	Accuracy	0.7955
5	Precision	0.7955
6	Recall	1.0000

```
[31]: print("Decision Tree Accuracy: \n",
          accuracy_score(y_test, y_predict_decisionTree))
```

```
print("Confusion Matrix: \n",
      confusion_matrix(y_test, y_predict_decisionTree))
```

Decision Tree Accuracy:

0.8275

Confusion Matrix:

```
[[ 491  327]
 [ 363 2819]]
```

```
[32]: cm = confusion_matrix(y_test, y_predict_decisionTree)
TN, FP, FN, TP = cm.ravel()
```

Create a Structured Table.

```
metrics_table = pd.DataFrame({
    'Metric': ['TP', 'TN', 'FP', 'FN', 'Accuracy', 'Precision', 'Recall'],
    'Value': [TP, TN, FP, FN,
              (TP+TN)/(TP+TN+FP+FN),
              TP/(TP+FP),
              TP/(TP+FN)]})
print(metrics_table)
```

	Metric	Value
0	TP	2819.000000
1	TN	491.000000
2	FP	327.000000
3	FN	363.000000
4	Accuracy	0.827500
5	Precision	0.896058
6	Recall	0.885921

Skills: Data cleaning and handling missing values. Exploratory Data Analysis (EDA). Binary classification using machine learning. Model evaluation using confusion and accuracy.

Task completed. Best wishes. Zaigham Abbas.